

# DM887 Assignment2 Q1

Jiawei Zhao

19.03.2024

---

**Algorithm 1** Least-Squares Temporal Differences (LSTD) Deep Q-Learning with Nonlinear Feature Extraction that maps a state to lower-dimensional latent embedding. While the action-value function should be a linear function of the output of the feature extractor.

---

- 1: Initialize a variational autoencoder  $\mathbf{A}$  as the feature extraction network with initial weights  $\theta_0$  [ $\theta_0$  are drawn from a zero-mean Gaussian distribution]
  - 2: Initialize the number of maximum episode  $N_0 = 10$ ,  $N_1 = 10$ , and  $N_2 = 30$  for each phase of training [It would be a wise choice to separate the warm-up phase, the autoencoder improvement phase, and the LSTD improvement phase] [Initialization of training hyperparameters]
  - 3: Initialize the number of maximum time step per episode  $t = 5000$
  - 4: Initialize the weights  $w_0$  for linear approximation of the action-value function  $y = f(\phi(s))$  which will be used for LSTD
  - 5: Initialize a replay memory buffer  $\mathcal{D}$  with a capacity  $N = t(N_0 + N_1 + N_2)$
  - 6: Initialize a discount factor  $\gamma = 0.9$  [A relatively large discount factor encourages long-term planning and faster convergence during training]
  
  - [The warm-up phase starts here]
  - 7: Freeze the weights of  $\mathbf{A}$  as  $\theta_0$
  - 8: Freeze the weights of the  $f(\phi(s))$  as  $w_0$
  - 9: **for** episode  $e = 1$  to  $N_0$  **do**
  - 10:     Initialize state  $s_0$
  - 11:     Preprocess  $s$  into  $s_1$  to adapt it as input of  $\mathbf{A}$  [preprocessing of high-dimensional states is necessary w.r.t. autoencoders]
  - 12:     **for** each time step  $t = 1$  to  $T$  **do**
  - 13:         Encode state  $s_t$  using  $\mathbf{A}$  to get latent embedding  $\phi(s_t)$
  - 14:         Select action  $a_t$  using  $\epsilon$ -greedy policy with  $\epsilon = 0.3$  [The learning curves of Game B of Assignment 1 corroborate that a relatively high  $\epsilon$  could improve  $Q$  more efficiently when an  $e$  does not end prematurely in the majority of cases]
  - 15:         Execute  $a_t$  and obtain reward  $r_t$  and new state  $s_{t+1}$
  - 16:         Store the transition of the current  $t$ , i.e.  $(s_t, a_t, r_t, s_{t+1})$  in  $\mathcal{D}$
  - 17:         Encode state  $s_t$  using  $\mathbf{A}$  to get latent embedding  $\phi(s_t)$
  - 18:         Select action  $a_t$  using  $\epsilon$ -greedy policy with  $\epsilon = 0.3$  [The learning curves of Game B of Assignment 1 corroborate that a relatively high  $\epsilon$  could improve  $Q$  more efficiently when an  $e$  does not end prematurely in the majority of cases]
  - 19:         Execute  $a_t$  and obtain reward  $r_t$  and new state  $s_{t+1}$
  - 20:         Store the transition of the current  $t$ , i.e.  $(s_t, a_t, r_t, s_{t+1})$  in  $\mathcal{D}$
  - 21:     **end for**
  - 22: **end for**
- [1]
-

---

**Algorithm 2** The LSTD update starts here

---

```
1: Unfreeze the weights of  $\mathbf{A}$ , i.e.  $\theta$ 
2: for episode  $e = 1$  to  $N_1$  do
3:   Repeat the same steps at phase 1
4:   for each time step  $t = 1$  to  $T$  do
5:     Repeat the same steps at phase 1
6:   [Start using the autoencoder algorithm to update  $w$ ]
7:     Sample a minibatch of transitions from  $\mathcal{D}$  with a batchsize with a randomly selected
8:   end for
9: end for
```

[1]

---

---

**Algorithm 3** The autoencoder + LSTD training starts here

---

```
1: Freeze the weights of  $\mathbf{A}$ , i.e.  $\theta$ 
2: Unfreeze the weights of  $f(\phi(s))$ 
3: for episode  $e = 1$  to  $N_2$  do
4:   Repeat the same steps at phase 1
5:   for each time step  $t = 1$  to  $T$  do
6:     Repeat the same steps at phase 1
7:     [Start using the online LSTD algorithm to update  $w$ ]
8:     Sample a minibatch of transitions from  $\mathcal{D}$  with a batchsize
9:   end for
10: end for
```

[1]

---